

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: pd.set_option("display.max_columns",None)
pd.set_option("display.max_rows",None)
```

Generating dataframe of train dataset

```
In [3]: df_train = pd.read_csv("D:\\titanic/train.csv")
df_train.head(2)
```

Out[3]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500	NaN	S
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599	71.2833	C85	C

Generating dataframe of test dataset

```
In [4]: df_test = pd.read_csv("D:\\titanic/test.csv")
df_test.head(2)
```

Out[4]:

	PassengerId	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
0	892	3	Kelly, Mr. James	male	34.5	0	0	330911	7.8292	NaN	Q
1	893	3	Wilkes, Mrs. James (Ellen Needs)	female	47.0	1	0	363272	7.0000	NaN	S

```
In [5]: df = pd.concat([df_train,df_test],axis=0)
df
```

14	15	0.0	0	Adolfina	female	14.00	0	0	330400	7.0042	NaN	S
15	16	1.0	2	Hewlett, Mrs. (Mary D Kingcome)	female	55.00	0	0	248706	16.0000	NaN	S
16	17	0.0	3	Rice, Master. Eugene	male	2.00	4	1	382652	29.1250	NaN	Q
17	18	1.0	2	Williams, Mr. Charles Eugene	male	NaN	0	0	244373	13.0000	NaN	S
18	19	0.0	3	Vander Planke, Mrs. Julius (Emelia Maria Vande...	female	31.00	1	0	345763	18.0000	NaN	S
19	20	1.0	3	Masselmani, Mrs. Fatima	female	NaN	0	0	2649	7.2250	NaN	C
20	21	0.0	2	Fynney, Mr. Joseph J	male	35.00	0	0	239865	26.0000	NaN	S
21	22	1.0	2	Beesley, Mr. Lawrence	male	34.00	0	0	248698	13.0000	D56	S
22	23	1.0	3	McGowan, Miss. Anna "Annie"	female	15.00	0	0	330923	8.0292	NaN	Q
23	24	1.0	1	Sloper, Mr. William Thompson	male	28.00	0	0	113788	35.5000	A6	S
24	25	0.0	3	Palsson, Miss. Torborg Danira	female	8.00	3	1	349909	21.0750	NaN	S

In [6]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 1309 entries, 0 to 417
Data columns (total 12 columns):
#   Column      Non-Null Count  Dtype
---  -
0   PassengerId  1309 non-null   int64
1   Survived     891 non-null    float64
2   Pclass       1309 non-null   int64
3   Name         1309 non-null   object
4   Sex          1309 non-null   object
5   Age         1046 non-null   float64
6   SibSp        1309 non-null   int64
7   Parch        1309 non-null   int64
8   Ticket       1309 non-null   object
9   Fare         1308 non-null   float64
10  Cabin        295 non-null    object
11  Embarked     1307 non-null   object
dtypes: float64(3), int64(4), object(5)
memory usage: 132.9+ KB
```

dropping those column which is not significant for the model preparation

```
In [7]: df_train.drop(['PassengerId', 'Cabin'], axis = 1, inplace=True)
df_train.head(10)
```

Out[7]:

	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Embarked
0	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500	S
1	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599	71.2833	C
2	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250	S
3	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000	S
4	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500	S
5	0	3	Moran, Mr. James	male	NaN	0	0	330877	8.4583	Q
6	0	1	McCarthy, Mr. Timothy J	male	54.0	0	0	17463	51.8625	S
7	0	3	Palsson, Master. Gosta Leonard	male	2.0	3	1	349909	21.0750	S
8	1	3	Johnson, Mrs. Oscar W (Elisabeth Vilhelmina Berg)	female	27.0	0	2	347742	11.1333	S
9	1	2	Nasser, Mrs. Nicholas (Adele Achem)	female	14.0	1	0	237736	30.0708	C

splitting name by respictive Title

```
In [8]: Name = df_train['Name']
df_train['Title'] = [ i.split('.')[0].split(',')[0] for i in Name]
df_train.head(10)
```

Out[8]:

	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Embarked	Title
0	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500	S	Mr
1	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599	71.2833	C	Mrs
2	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250	S	Miss
3	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000	S	Mrs
4	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500	S	Mr
5	0	3	Moran, Mr. James	male	NaN	0	0	330877	8.4583	Q	Mr
6	0	1	McCarthy, Mr. Timothy J	male	54.0	0	0	17463	51.8625	S	Mr
7	0	3	Palsson, Master. Gosta Leonard	male	2.0	3	1	349909	21.0750	S	Master
8	1	3	Johnson, Mrs. Oscar W (Elisabeth Vilhelmina Berg)	female	27.0	0	2	347742	11.1333	S	Mrs
9	1	2	Nasser, Mrs. Nicholas (Adele Achem)	female	14.0	1	0	237736	30.0708	C	Mrs

```
In [9]: df_train.drop("Name",axis=1,inplace=True)
df_train.head(10)
```

Out[9]:

	Survived	Pclass	Sex	Age	SibSp	Parch	Ticket	Fare	Embarked	Title
0	0	3	male	22.0	1	0	A/5 21171	7.2500	S	Mr
1	1	1	female	38.0	1	0	PC 17599	71.2833	C	Mrs
2	1	3	female	26.0	0	0	STON/O2. 3101282	7.9250	S	Miss
3	1	1	female	35.0	1	0	113803	53.1000	S	Mrs
4	0	3	male	35.0	0	0	373450	8.0500	S	Mr
5	0	3	male	NaN	0	0	330877	8.4583	Q	Mr
6	0	1	male	54.0	0	0	17463	51.8625	S	Mr
7	0	3	male	2.0	3	1	349909	21.0750	S	Master
8	1	3	female	27.0	0	2	347742	11.1333	S	Mrs
9	1	2	female	14.0	1	0	237736	30.0708	C	Mrs

define a class for data cleaning

```
In [10]: def data_cleaning(data):
          d = df_train[data]
          for column in d.columns:
              if df_train[column].isnull().any():
                  df_train[column].fillna(df_train[column].median(),inplace = True)
```

```
In [11]: feature = ['Age', 'SibSp', 'Parch', 'Fare']
data_cleaning(feature)
```

```
In [12]: df_train['Embarked'].fillna(method='ffill',inplace=True)
```


```
In [13]: print(df_train.isnull().sum())
```


```
Survived    0  
Pclass      0  
Sex         0  
Age         0  
SibSp       0  
Parch       0  
Ticket      0  
Fare        0  
Embarked    0  
Title       0  
dtype: int64
```

univariate analysis for feature having numerical datatype

```
In [14]: def plot(feature):  
         v = df_train[feature]  
         for i in v.columns:  
             sns.displot(df_train[i],kde=True)  
             plt.show()
```

```
In [15]: feature = ['Age', 'Fare']  
plot(feature)
```






univariate analysis for feature having categorical datatypes

```
In [16]: def plot(feature):  
         v = df_train[feature]  
         v_value = v.value_counts()  
         plt.figure(figsize = (9,6))  
         plt.xticks(rotation = 60)  
         plt.bar(v_value.index,v_value)  
         plt.show()
```


```
In [17]: feature = ['Pclass', 'Sex', 'SibSp', 'Parch', 'Embarked', 'Title']  
         for i in feature:  
             plot(i)
```



Multivariate Analysis

```
In [18]: sns.pairplot(df_train,hue="Survived")  
plt.show()
```





```
In [19]: df_train.head(10)
```


```
Out[19]:
```

	Survived	Pclass	Sex	Age	SibSp	Parch	Ticket	Fare	Embarked	Title
0	0	3	male	22.0	1	0	A/5 21171	7.2500	S	Mr
1	1	1	female	38.0	1	0	PC 17599	71.2833	C	Mrs
2	1	3	female	26.0	0	0	STON/O2. 3101282	7.9250	S	Miss
3	1	1	female	35.0	1	0	113803	53.1000	S	Mrs
4	0	3	male	35.0	0	0	373450	8.0500	S	Mr
5	0	3	male	28.0	0	0	330877	8.4583	Q	Mr
6	0	1	male	54.0	0	0	17463	51.8625	S	Mr
7	0	3	male	2.0	3	1	349909	21.0750	S	Master
8	1	3	female	27.0	0	2	347742	11.1333	S	Mrs
9	1	2	female	14.0	1	0	237736	30.0708	C	Mrs

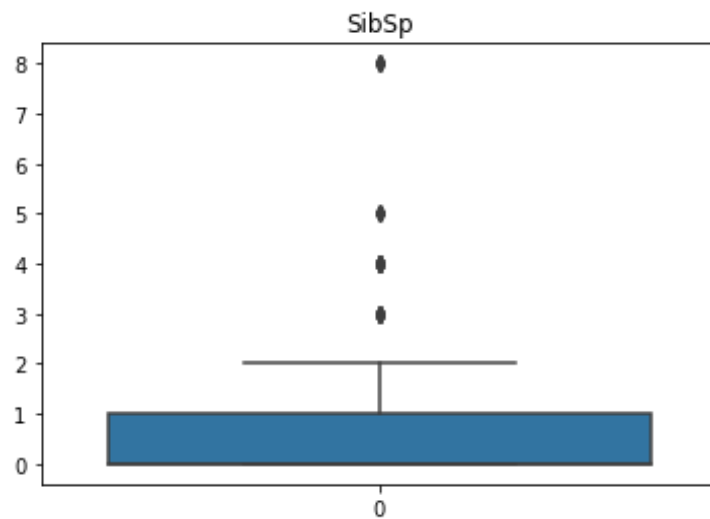
Finding Outliers

```
In [20]: def detect_outliers(feature):
threshold = 3
outliers = []
data = df_train[feature]
mean = np.mean(data)
std = np.std(data)
for x in data:
    z_score = ( x - mean ) / std
    if z_score > threshold :
        outliers.append(x)
plt.title(f"{i}")
sns.boxplot(data=data)
plt.show()
print(f"outliers of {feature} column:- \n {outliers}")
```

```
In [21]: features = ["Age", "SibSp", "Parch", "Fare"]  
for i in features :  
    detect_outliers(i)
```




outliers of Age column:-
[71.0, 70.5, 71.0, 80.0, 70.0, 70.0, 74.0]




outliers of SibSp column:-

[4, 4, 5, 4, 5, 4, 8, 4, 4, 8, 4, 8, 4, 4, 4, 4, 8, 5, 5, 4, 4, 5, 4, 4, 8, 4, 4, 8, 4, 8]



outliers of Parch column:-

[5, 5, 3, 4, 4, 3, 4, 4, 5, 5, 6, 3, 3, 3, 5]



outliers of Fare column:-

[263.0, 263.0, 247.5208, 512.3292, 247.5208, 262.375, 263.0, 211.5, 227.525, 263.0, 221.7792, 227.525, 512.3292, 211.3375, 227.525, 227.525, 211.3375, 512.3292, 262.375, 211.3375]

```
In [22]: df_train.head(10)
```

```
Out[22]:
```

	Survived	Pclass	Sex	Age	SibSp	Parch	Ticket	Fare	Embarked	Title
0	0	3	male	22.0	1	0	A/5 21171	7.2500	S	Mr
1	1	1	female	38.0	1	0	PC 17599	71.2833	C	Mrs
2	1	3	female	26.0	0	0	STON/O2. 3101282	7.9250	S	Miss
3	1	1	female	35.0	1	0	113803	53.1000	S	Mrs
4	0	3	male	35.0	0	0	373450	8.0500	S	Mr
5	0	3	male	28.0	0	0	330877	8.4583	Q	Mr
6	0	1	male	54.0	0	0	17463	51.8625	S	Mr
7	0	3	male	2.0	3	1	349909	21.0750	S	Master
8	1	3	female	27.0	0	2	347742	11.1333	S	Mrs
9	1	2	female	14.0	1	0	237736	30.0708	C	Mrs

```
In [23]: #outliers of Age column:- [71.0, 70.5, 71.0, 80.0, 70.0, 70.0, 74.0]
o = [71.0, 70.5, 71.0, 80.0, 70.0, 70.0, 74.0]
a = df_train['Age']
for i in a:
    if i in o:
        df_train.Age.replace(df_train['Age'].median(),inplace=True)
```

```
In [24]: # outliers of Parch column:- [5, 5, 3, 4, 4, 3, 4, 4, 5, 5, 6, 3, 3, 3, 5]
o = [5, 5, 3, 4, 4, 3, 4, 4, 5, 5, 6, 3, 3, 3, 5]
a = df_train['Parch']
for i in a:
    if i in o:
        df_train.Parch.replace(df_train['Parch'].median(),inplace=True)
```



```
In [25]: # outliers of Fare column:- [263.0, 263.0, 247.5208, 512.3292, 247.5208, 262.375, 263.0, 211.5, 227.525, 263.0, 221.779]
o = [263.0, 263.0, 247.5208, 512.3292, 247.5208, 262.375, 263.0, 211.5, 227.525, 263.0, 221.7792, 227.525, 512.3292, 211.5]
a = df_train['Fare']
for i in a:
    if i in o:
        df_train.Fare.replace(df_train['Fare'].median(),inplace=True)
```


```
In [26]: # outliers of SibSp column:- [4, 4, 5, 4, 5, 4, 8, 4, 4, 8, 4, 8, 4, 4, 4, 4, 8, 5, 5, 4, 4, 5, 4, 4, 8, 4, 4, 8, 4, 8]
o = [4, 4, 5, 4, 5, 4, 8, 4, 4, 8, 4, 8, 4, 4, 4, 4, 8, 5, 5, 4, 4, 5, 4, 4, 8, 4, 4, 8, 4, 8]
for i in a:
    if i in o:
        df_train.SibSp.replace(df_train['SibSp'].median(),inplace=True)
```

```
In [27]: df_train.head(10)
```

Out[27]:

	Survived	Pclass	Sex	Age	SibSp	Parch	Ticket	Fare	Embarked	Title
0	0	3	male	22.0	1	0	A/5 21171	7.2500	S	Mr
1	1	1	female	38.0	1	0	PC 17599	71.2833	C	Mrs
2	1	3	female	38.0	1	0	STON/O2. 3101282	7.9250	S	Miss
3	1	1	female	35.0	1	0	113803	53.1000	S	Mrs
4	0	3	male	35.0	1	0	373450	8.0500	S	Mr
5	0	3	male	35.0	1	0	330877	8.4583	Q	Mr
6	0	1	male	54.0	1	0	17463	51.8625	S	Mr
7	0	3	male	2.0	3	0	349909	21.0750	S	Master
8	1	3	female	2.0	3	0	347742	11.1333	S	Mrs
9	1	2	female	14.0	1	0	237736	30.0708	C	Mrs

```
In [28]: l = ["SibSp", "Parch", "Survived", "Age", "Fare"]
sns.heatmap(df[l].corr(), annot = True, fmt = ".2f")
plt.show()
```




```
In [29]: df_train.head()
```

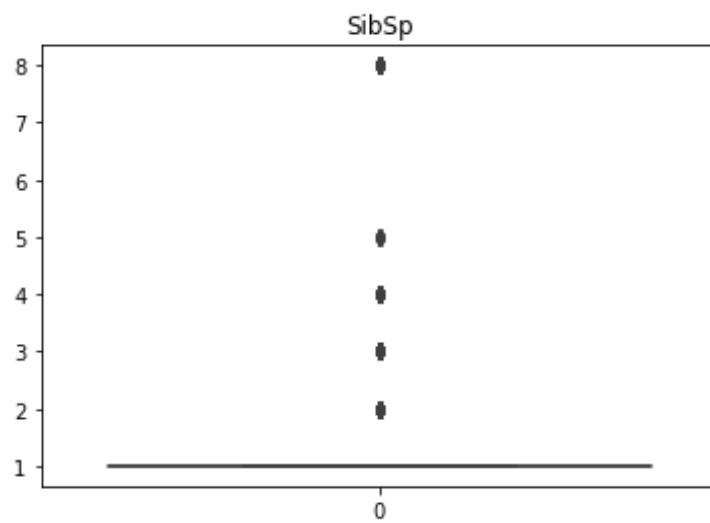
Out[29]:

	Survived	Pclass	Sex	Age	SibSp	Parch	Ticket	Fare	Embarked	Title
0	0	3	male	22.0	1	0	A/5 21171	7.2500	S	Mr
1	1	1	female	38.0	1	0	PC 17599	71.2833	C	Mrs
2	1	3	female	38.0	1	0	STON/O2. 3101282	7.9250	S	Miss
3	1	1	female	35.0	1	0	113803	53.1000	S	Mrs
4	0	3	male	35.0	1	0	373450	8.0500	S	Mr

```
In [30]: for i in features :  
         detect_outliers(i)
```

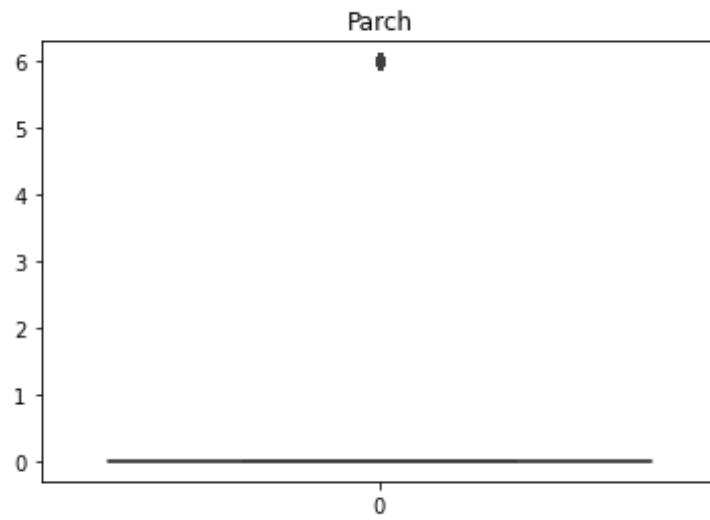


outliers of Age column:-
[80.0]




outliers of SibSp column:-

```
[8, 8, 8, 8, 8, 8, 8, 8, 8, 8, 8, 8, 8, 8, 8, 8, 8, 8, 8, 8, 8, 8, 8, 8, 8, 8]
```



outliers of Parch column:-

```
[]
```



outliers of Fare column:-

[263.0, 263.0, 247.5208, 512.3292, 247.5208, 262.375, 263.0, 263.0, 263.0, 263.0, 263.0, 263.0, 211.5, 227.525, 263.0, 221.7792, 227.525, 512.3292, 211.3375, 227.525, 227.525, 211.3375, 512.3292, 262.375, 211.3375]

```
In [32]: df_train.head(10)
```

Out[32]:

	Survived	Pclass	Sex	Age	SibSp	Parch	Ticket	Fare	Embarked	Title
0	0	3	male	22.0	1	0	A/5 21171	7.2500	S	Mr
1	1	1	female	38.0	1	0	PC 17599	71.2833	C	Mrs
2	1	3	female	38.0	1	0	STON/O2. 3101282	7.9250	S	Miss
3	1	1	female	35.0	1	0	113803	53.1000	S	Mrs
4	0	3	male	35.0	1	0	373450	8.0500	S	Mr
5	0	3	male	35.0	1	0	330877	8.4583	Q	Mr
6	0	1	male	54.0	1	0	17463	51.8625	S	Mr
7	0	3	male	2.0	3	0	349909	21.0750	S	Master
8	1	3	female	2.0	3	0	347742	11.1333	S	Mrs
9	1	2	female	14.0	1	0	237736	30.0708	C	Mrs

```
In [33]: df_train=pd.get_dummies(df_train,columns=["Title"])
df_train = pd.get_dummies(df_train, columns=["Embarked"])
df_train = pd.get_dummies(df_train, columns= ["Ticket"], prefix = "T")
df_train["Sex"] = df_train["Sex"].astype("category")
df_train = pd.get_dummies(df_train, columns=["Sex"])
df_train.head(2)
```

Out[33]:

	Survived	Pclass	Age	SibSp	Parch	Fare	Title_	Title_	Title_	Title_	Title_	Title_	Title_	Title_	Title_	Title_	Title_	Title_	Title_
							Capt	Col	Don	Dr	Jonkheer	Lady	Major	Master	Miss	Mlle	Mme	Mr	Mrs
0	0	3	22.0	1	0	7.2500	0	0	0	0	0	0	0	0	0	0	0	1	0
1	1	1	38.0	1	0	71.2833	0	0	0	0	0	0	0	0	0	0	0	0	1

```
In [34]: df_train.head(10)
```

Out[34]:

	Survived	Pclass	Age	SibSp	Parch	Fare	Title_ Capt	Title_ Col	Title_ Don	Title_ Dr	Title_ Jonkheer	Title_ Lady	Title_ Major	Title_ Master	Title_ Miss	Title_ Mlle	Title_ Mme	Title_ Mr	Title_ Mrs
0	0	3	22.0	1	0	7.2500	0	0	0	0	0	0	0	0	0	0	0	1	0
1	1	1	38.0	1	0	71.2833	0	0	0	0	0	0	0	0	0	0	0	0	1
2	1	3	38.0	1	0	7.9250	0	0	0	0	0	0	0	0	1	0	0	0	0
3	1	1	35.0	1	0	53.1000	0	0	0	0	0	0	0	0	0	0	0	0	1
4	0	3	35.0	1	0	8.0500	0	0	0	0	0	0	0	0	0	0	0	1	0
5	0	3	35.0	1	0	8.4583	0	0	0	0	0	0	0	0	0	0	0	1	0
6	0	1	54.0	1	0	51.8625	0	0	0	0	0	0	0	0	0	0	0	1	0
7	0	3	2.0	3	0	21.0750	0	0	0	0	0	0	0	1	0	0	0	0	0
8	1	3	2.0	3	0	11.1333	0	0	0	0	0	0	0	0	0	0	0	0	1
9	1	2	14.0	1	0	30.0708	0	0	0	0	0	0	0	0	0	0	0	0	1

```
In [50]: from sklearn.model_selection import train_test_split, StratifiedKFold, GridSearchCV
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier, VotingClassifier, GradientBoostingClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score
```

```
In [51]: x = df_train.drop("Survived", axis = 1)
y = df_train["Survived"]
x_train, x_test, y_train, y_test = train_test_split(x,y, test_size = 0.33, random_state = 42)
```

```
In [52]: model = LogisticRegression()  
model.fit(x_train,y_train)
```

C:\Users\kants\AppData\Local\Programs\Python\Python37\lib\site-packages\sklearn\linear_model_logistic.py:765: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)

```
Out[52]: LogisticRegression()
```

```
In [53]: y_predicted=model.predict(x_test)
```

```
In [54]: model.score(x_test,y_test)
```

```
Out[54]: 0.823728813559322
```

```
In [55]: from sklearn.metrics import r2_score,mean_squared_error  
print(f'R^2 : {r2_score(y_test,y_predicted)}')  
print(f'MSE : {mean_squared_error(y_test,y_predicted)}')  
print(f'RMSE: {np.sqrt(mean_squared_error(y_test, y_predicted))}')  

```

R^2 : 0.2695238095238095

MSE : 0.17627118644067796

RMSE: 0.41984662251907895


```
In [56]: random_state = 42
classifier = [DecisionTreeClassifier(random_state = random_state),
              SVC(random_state = random_state),
              RandomForestClassifier(random_state = random_state),
              LogisticRegression(random_state = random_state),
              KNeighborsClassifier()]

dt_grid_parameter = {"min_samples_split" : range(10,500,20),
                     "max_depth": range(1,20,2)}

svc_grid_parameter = {"kernel" : ["rbf"],
                      "gamma": [0.001, 0.01, 0.1, 1],
                      "C": [1,10,50,100,200,300,1000]}

rf_grid_parameter = {"max_features": [1,3,10],
                     "min_samples_split": [2,3,10],
                     "min_samples_leaf": [1,3,10],
                     "bootstrap": [False],
                     "n_estimators": [100,300],
                     "criterion": ["gini"]}

lr_grid_parameter = {"C": np.logspace(-3,3,7),
                     "penalty": ["l1","l2"]}

knn_grid_parameter = {"n_neighbors": np.linspace(1,19,10, dtype = int).tolist(),
                      "weights": ["uniform","distance"],
                      "metric": ["euclidean","manhattan"]}

classifier_param = [dt_grid_parameter,
                    svc_grid_parameter,
                    rf_grid_parameter,
                    lr_grid_parameter,
                    knn_grid_parameter]
```

```
In [58]: cv_result = []
best_estimators = []
for i in range(len(classifier)):
    clf = GridSearchCV(classifier[i], param_grid=classifier_param[i], cv = StratifiedKFold(n_splits = 10), scoring = "a
    clf.fit(x_train,y_train)
    cv_result.append(clf.best_score_)
    best_estimators.append(clf.best_estimator_)
    print(cv_result[i])
```

Fitting 10 folds for each of 250 candidates, totalling 2500 fits

0.8238983050847457

Fitting 10 folds for each of 28 candidates, totalling 280 fits

0.7734745762711865

Fitting 10 folds for each of 54 candidates, totalling 540 fits

0.8389830508474576

Fitting 10 folds for each of 14 candidates, totalling 140 fits

C:\Users\kants\AppData\Local\Programs\Python\Python37\lib\site-packages\sklearn\model_selection_search.py:925: UserWarning: One or more of the test scores are non-finite: [nan 0.68446328 nan 0.78163842 nan 0.7919774

nan 0.80706215 nan 0.78183616 nan 0.78019774

nan 0.8070339]

category=UserWarning

C:\Users\kants\AppData\Local\Programs\Python\Python37\lib\site-packages\sklearn\linear_model_logistic.py:765: ConvergenceWarning: lbfgs failed to converge (status=1):

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)

0.8070621468926553

Fitting 10 folds for each of 40 candidates, totalling 400 fits

0.7315536723163841


```
In [59]: cv_results = pd.DataFrame({"Cross Validation Means":cv_result, "ML Models":["DecisionTreeClassifier", "SVM","RandomFore",
                                         "LogisticRegression",
                                         "KNeighborsClassifier"]})

g = sns.barplot("Cross Validation Means", "ML Models", data = cv_results)
g.set_xlabel("Mean Accuracy")
g.set_title("Cross Validation Scores")
```

C:\Users\kants\AppData\Local\Programs\Python\Python37\lib\site-packages\seaborn_decorators.py:43: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

FutureWarning

Out[59]: Text(0.5, 1.0, 'Cross Validation Scores')



```
In [63]: poll = VotingClassifier(estimators = [("dt",best_estimators[0]),
                                             ("rfc",best_estimators[2]),
                                             ("lr",best_estimators[3])],
                                voting = "soft", n_jobs = -1)

poll = votingC.fit(x_train, y_train)
print(accuracy_score(votingC.predict(x_test),y_test))
```

0.8305084745762712

In []: